

CCX: Composable Custom eXtensions

2023 RISC-V Summit Europe (Barcelona)

https://www.youtube.com/watch?v=ugr3RqmWhWw&list=PL85jopFZCnbMIDT5WOagUjQnk1ZSgjuRa&index=10&ab_channel=RISC-VInternational

Guy Lemieux (UBC) and Jan Gray (Gray Research)

Other Contributors

Tim Vogt (Lattice), Tim Callahan (Google), Charles Papon (SpinalHDL),
Maciej Kurc (Antmicro), Karol Gugala (Antmicro), Olof Kindren (Qamcom)

The problem: incompatible interfaces






RISC-V
Processor



RISC-V
Accelerator

The problem (2): the stack



Host		Peripheral	
Target Product (example)	USB F/W	USB F/W	Target Product (example)
	Mass Storage Class Driver (HMSC)	Mass Storage Class Driver (PMSC)	 USB Thumb Drive
	Communication Class Driver (HCDC)	Communication Class Driver (PCDC)	 RS232C-USB Conversion

- Standard physical interface
- Needs standardized software stack

Common Device Class GUIDs

Class	GUID	Device Description
CDROM	{4D36E965-E325-11CE-BFC1-08002BE10318}	CD/DVD/Blu-ray drives
DiskDrive	{4D36E967-E325-11CE-BFC1-08002BE10318}	Hard drives
Display	{4D36E968-E325-11CE-BFC1-08002BE10318}	Video adapters
FDC	{4D36E969-E325-11CE-BFC1-08002BE10318}	Floppy controllers
FloppyDisk	{4D36E980-E325-11CE-BFC1-08002BE10318}	Floppy drives
HDC	{4D36E96A-E325-11CE-BFC1-08002BE10318}	Hard drive controllers
HIDClass	{745A17A0-74D3-11D0-B6FE-00A0C90F57DA}	Some USB devices
1394	{6BDD1FC1-810F-11D0-BEC7-08002BE2092F}	IEEE 1394 host controller

The problem (3): scarce ISA encodings

- 32b opcodes → nearly out of space
 - Even RV32 **Standard Extension** encodings **will overlap**

*“the **P** and **V** extensions have overlapping functionality, yet are largely suited to different application domains, we suggest declaring them mutually incompatible.*


This avails each extension of the other's major opcode” (ARC)

- Problem is worse for accelerators
 - unlike Standard Extensions, custom instructions are unmanaged

Many accelerator interfaces

- RoCC (UC Berkeley Rocket)
- Core-V-XIF (OpenHW Group)
- SCAIE-V (TU Darmstadt)
- ACE (Andes)
- VCIX (SiFive for vectors)

...more...



Common ground?

These interfaces solve the same problem:

Define new custom ISA encodings

... to connect a few instruction-based Accelerators

... to a few CPUs

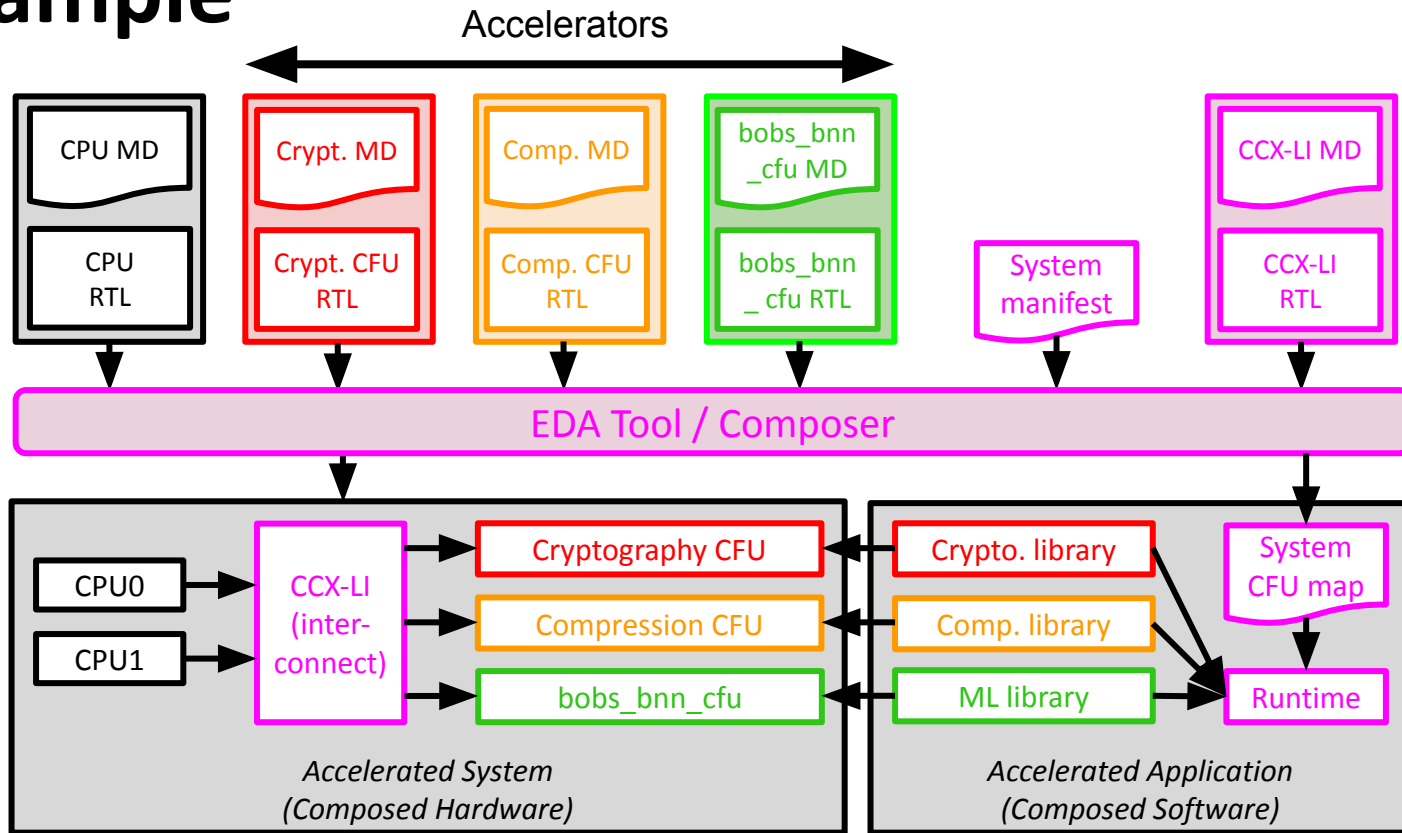
... to run software.

Why is this not enough?

Because your extensions won't compose with my extensions.

This continuously fragments the ecosystem!

Example



A programming model


1. Discover
2. Select and allocate context (eg, constructor)
3. Issue custom instructions
4. Deselect (eg, destructor)

```
if (CI bm(GUID_BitmanipCI); bm)           // csrrw x2,mcfu_selector,x1
    count = custom0( POPCNT, data, 0 );    // custom-0 POPCNT,x3,x4,x5
    ~CI() deselection implied              // csrw mcfu_selector,x2
else
    count = popcount(data);                // unaccelerated software
```


Features

- Composability
 - Resolve ISA encoding conflicts
 - No centralized management
- Versioning & run-time discovery
- Multiple state contexts (& OS protected)
- Hierarchical nesting (figure on next slide)


Hierarchical accelerator composition



Suggestions: (1) consider merging cfu_reg and cfu_imm into custom-0, (2) consider “official” way of how to add official RISC-V extensions (eg, RVV) as a CCX implementation.

What's needed?

- GUID (128b) – self-minted for each version of accelerator
- 2 CSRs
 - “Selector” to choose CFU_ID[7:0], STATE_ID[7:0]
 - “Status” to accrue errors



- Standardized instruction encodings
 - (custom-0 block) cfu_reg cf_id[10],rd,rs1,rs2
 - (custom-1 block) cfu_imm cf_id[7],rd,rs1,imm8
 - (custom-2 block) cfu_flex cf_id[25] // also sends [rs1], [rs2]

A programming model

1. Discover
2. Select and allocate context (eg, constructor)
3. Issue custom instructions
4. Deselect (eg, destructor)

```
if (CI bm(GUID_BitmanipCI); bm)           // csrrw x2,mcfu_selector,x1
    count = custom0( POPCNT, data, 0 );    // custom-0 POPCNT,x3,x4,x5
    ~CI() deselection implied              // csr w mcfu_selector,x2
else
    count = popcount(data);                // unaccelerated software
```

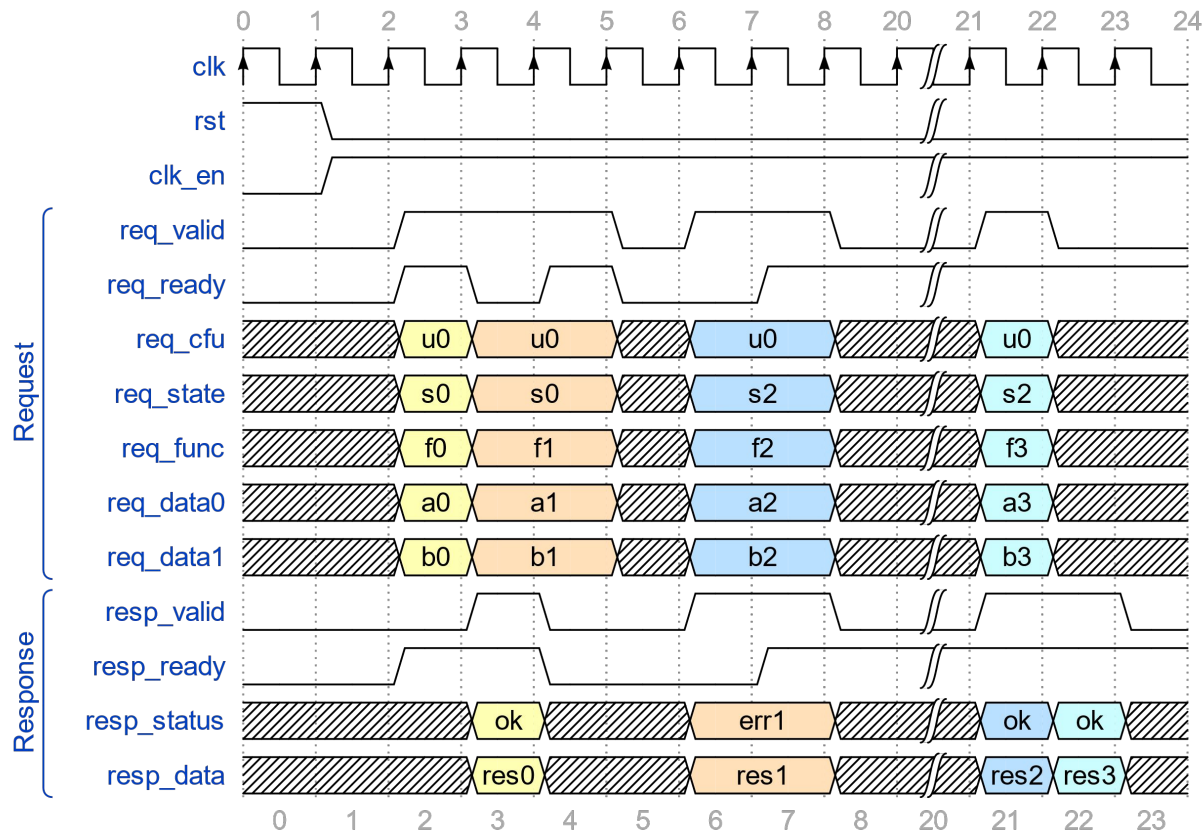
Privileged mode

Privileged mode needs 2 additional CSRs

- Protected CSR
 - CFU Selector Table
 - pointer to 4KiB page in DRAM (1024 entries)
 - each entry has **< CFU index, STATE index >**
- User CSR
 - offset to one of 1024 entries in CFU Selector Table

Physical Logic Interface

- Compatible with AXI-Stream



What else?

- Software

- four custom-0 instructions to manage state
 - `cf_read_status`, `cf_read_state`
 - `cf_write_status`, `cf_write_state`
- run-time system for discovery, allocation, context switching

- Chip building

- **System manifest** describes CPU + accelerator connections
- **Composer (EDA tool)** to connect CPUs + accelerators

Status



To Do

- Implementation proofs
 - CFU runtime
 - Composition tool
 - RTL/FPGA proofs (plugfest)
- Wishlist
 - Speculative (issue+squash)
 - Virtual memory
 - Coherence + consistency
 - Hypervisors

Call to Action

We propose two new Task Groups:

CCX-ISA: ISA and CSRs to enable composable custom extensions

CCX-LI: logic interface to physically connect accelerators with CPUs

To bootstrap this, we respectfully submit:



*Draft Proposed RISC-V Composable Custom
Extensions Specification*

<https://github.com/grayresearch/CFU>